

AMANDA COSTON

Email: amandacoston@microsoft.com
Phone: 703-401-1212

Employment

2023-2024 MICROSOFT RESEARCH (MSR) NEW ENGLAND
Postdoc researcher, Machine Learning and Statistics

2024- UNIVERSITY OF CALIFORNIA, BERKELEY
Assistant Professor, Department of Statistics

Education

2017-2023 CARNEGIE MELLON UNIVERSITY
Ph.D. candidate in Machine Learning and Public Policy
Advisors: [Alexandra Chouldechova](#) & [Edward Kennedy](#)
Thesis: “Principled machine learning for societally consequential decision making”.
Committee: [Edward Kennedy](#), [Alexandra Chouldechova](#), [Hoda Heidari](#), & [Sendhil Mullainathan](#)

2017-2019 CARNEGIE MELLON UNIVERSITY
M.S. in Machine Learning.

2009-2013 PRINCETON UNIVERSITY
B.S.E. *magna cum laude* in Computer Science
Certificate in the Princeton School of Public and International Affairs
Advisor: [Robert Schapire](#)
Thesis: “Machine learning techniques for the diagnosis of pediatric tuberculosis”.

Selected Awards & Honors

Research

2023 Best Paper Award at [ACM Conference on Fairness, Accountability, and Transparency \(FAccT\)](#)

2023 William W. Cooper Doctoral Dissertation Award

2023 Best Paper Award at [IEEE Conference on Secure and Trustworthy Machine Learning \(SaTML\)](#)

2022 [Rising Star in EECS](#) by UT-Austin

2022 [Rising Star in Machine Learning](#) by University of Maryland

2022 [Rising Star in Data Science](#) by University of Chicago

2022 [Meta Research PhD Fellow](#)

2022 Future Leader in Responsible Data Science by University of Michigan Institute for Data Science

2020 K&L Gates Presidential Fellow in Ethics and Computational Technologies

2019 NSF Graduate Research Fellow

2019 Tata Consultancy Services Presidential Fellow

2019 Suresh Konda Best First Paper Award by Heinz College of Carnegie Mellon University

Service

2020 Carolyn Comer Graduate Student Involvement Award by Carnegie Mellon University

2013 Department of Computer Science Service Award by Princeton University

Research & Industry Experience

- 2021 FACEBOOK AI APPLIED RESEARCH (FAIAR)
Research intern, Responsible AI
Conducted a creator-centric fairness assessment of Instagram Reels.
- 2020 REGLAB, STANFORD UNIVERSITY
Research Fellow, Regulation, Evaluation, and Governance Lab at Stanford Law School
Conducted audit of mobility data for racial bias.
- 2018 IBM RESEARCH AI
Science for Social Good Fellow
Developed methods for fairness-aware learning under domain shift.
- 2017 HIVISASA
Technical Consultant, Kenya
Built full-stack analytics for citizen journalism website.
- 2015-2017 TENEO
Data Scientist
- 2013-2015 MICROSOFT
Program Manager, Bing
- 2010-2011 SHELTON PSYCHOLOGY LAB, PRINCETON UNIVERSITY
Research Assistant
Administered experiments testing stereotype priming effect on STEM performance

Research Interests

Theory: causal inference, machine learning, algorithmic fairness & societal impacts
Application: child welfare, consumer credit lending, criminal justice, health policy

Publications & Manuscripts

* indicates joint lead authors

- Working Papers**
- Coston A, Kennedy EH. Counterfactual audit of racial bias in police traffic stops. *American Causal Inference Conference (ACIC) 2022* oral presentation (20% selection rate).
- Coston A, Kennedy EH. The role of the geometric mean in case-control studies. [arxiv.org:2207.09016](https://arxiv.org/abs/2207.09016)
- Rambachan A, Coston A, Kennedy EH. Robust design and evaluation of predictive algorithms under unmeasured confounding. *ACIC 2022. NeurIPS 2022 Workshop on Algorithmic Fairness through the Lens of Causality and Privacy*. [arxiv.org:2212.09844](https://arxiv.org/abs/2212.09844)
- Guerdan L, Coston A, Wu ZS, Holstein K. Subgroup policy auditing under confounding.

Kawakami A, Coston A, Heidari H, Y Zhu, Holstein K. Situating AI: Co-designing assessments to forefront context-specific needs for AI design and deployment.

Kawakami A, Coston A, Heidari H, Holstein K, Zhu Y. Studying up public sector AI: Shifting our gaze upwards to study systems of power in public sector AI. *ACM conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO 2023)*. oral presentation (18% selection rate)

Publications

Guerdan L, Coston A, Wu ZS, Holstein K. Counterfactual decision support under outcome measurement error. *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. 2023; 1584–1598. doi:10.1145/3593013.3594101. (arxiv.org:2302.11121) **Best Paper Award** by FAccT

Field A, Coston A, Gandhi N, Chouldechova A, Putnam-Hornstein E, Steier D, Tsvetkov Y. Examining risks of racial biases in NLP tools for Child Protective Services. *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. 2023; 1479–1492. doi:10.1145/3593013.3594094

Guerdan L, Coston A, Wu ZS, Holstein K. Ground(Less) truth: A causal framework for proxy labels in human-algorithm decision making. *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. 2023; 688–704. doi:10.1145/3593013.3594036 (arxiv.org:2302.06503)

Coston A, Kawakami A, Zhu Y, Holstein K, Heidari H. A validity perspective on evaluating the justified use of data-driven decision-making algorithms. *IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*. 2023. (arxiv.org:2206.14983). **Best Paper Award** by SaTML

Coston A*, Rambachan A*, Chouldechova A. Characterizing fairness over the set of good models under selective labels. *International Conference on Machine Learning 139 (ICML)*. 2021; 2144–2155. <http://proceedings.mlr.press/...> (arxiv.org:2101.00352)

Coston A, Guha N, Ouyang D, Lu L, Chouldechova A, Ho DE. Leveraging administrative data for bias audits: Assessing disparate coverage with mobility data for COVID-19 policy. *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. 2021; 173–184. doi:10.1145/3442188.3445881 (arxiv.org:2011.07194)

Coston A, Kennedy EH, Chouldechova A. Counterfactual predictions under runtime confounding. *Advances in Neural Information Processing Systems 33 (NeurIPS)*. 2020; 4150–4162. <https://papers.nips.cc/paper/...> (arxiv.org:2006.16916)

Coston A, Mishler A, Kennedy EH, Chouldechova A. Counterfactual risk assessments, evaluation, and fairness. *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. 2020; 582–593. doi:10.1145/3351095.3372851 (arxiv.org:1909.00066)

Zhao H, Coston A, Adel T, Gordon GJ. Conditional learning of fair representations. *International Conference on Learning Representations (ICLR)*. 2020. <https://iclr.cc/...> (arxiv.org:1910.07162)

Li L, Zuo R, Coston A, Weiss JC, Chen GH. Neural topic models with survival supervision: Jointly predicting time-to-event outcomes and learning how clinical features relate. *International Conference on Artificial Intelligence in Medicine (AIME)*. 2020; 371-381. <https://link.springer.com/...> (arxiv.org:2007.07796)

Coston A, Ramamurthy KN, Wei D, Varshney KR, Speakman S, Mustahsan Z, Chakraborty S. Fair transfer learning with missing protected attributes. *Proceedings of the AAAI / ACM Conference on Artificial Intelligence, Ethics, and Society (AIES)*. 2019; 91-98. doi:10.1145/3306618.3314236

**Book
Chapter**

Coston A, Rubio MD, Kennedy EH. Statistical analysis of randomized experiments. *AI for Social Impact*. ai4socialbook.org

**Peer-reviewed
non-archival
papers**

Kawakami A, Coston A, Heidari H, Holstein K, Zhu Y. Studying up public sector AI: Shifting our gaze upwards to study systems of power in public sector AI. *ACM conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO 2023)*. oral presentation (18% selection rate)

Kawakami A, Coston A, Y Zhu, Heidari H, Holstein K. Recentering validity considerations through early-stage deliberations around AI and policy design. *CHI 2023 Workshop on Designing Technology and Policy*.

Rambachan A, Coston A, Kennedy EH. Counterfactual risk assessments under unmeasured confounding. *ACIC 2022. NeurIPS 2022 Workshop on Algorithmic Fairness through the Lens of Causality and Privacy*. [arxiv.org:2212.09844](https://arxiv.org/abs/2212.09844)

Guerdan L, Coston A, Wu ZS, Holstein K. Counterfactual decision support under treatment-conditional outcome measurement error. *NeurIPS 2022 Workshop on Causality for Real-world Impact*.

Guerdan L, Coston A, Wu ZS, Holstein K. Ground(less) truth: The problem with proxy outcomes in human-AI decision making. *NeurIPS 2022 Workshop on Human-Centered AI*.

Coston A, Kennedy EH. Counterfactual audit of racial bias in police traffic stops. *ACIC 2022* oral presentation (20% selection rate).

Coston A, Kawakami A, Zhu Y, Holstein K, Heidari H. A validity perspective on evaluating the justified use of data-driven decision-making algorithms. *ACM conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO 2022)*. [arxiv.org:2206.14983](https://arxiv.org/abs/2206.14983)

Coston A, Kennedy EH, Chouldechova A. Counterfactual risk assessments, evaluation, and fairness. *NeurIPS 2019 Workshop on Causal Machine Learning*.

Coston A, Kennedy EH, Chouldechova A. Counterfactual risk assessments and evaluation for child welfare screening. *ACIC 2019*.

Coston A, Leqi L. Offline heterogeneous policy evaluation: A causal approach. *ICML 2018 Workshop on Causal ML*.

Presentations

** indicates presentation scheduled for future date*

Invited Talks

2024*	Applied Statistics Workshop, Department of Government and Institute for Quantitative Social Science, Harvard University, Boston, MA
2023*	Bringing Statistical Thinking into Fair, Transferrable Machine Learning, Institute of Mathematical Sciences International Conference on Statistics and Data Science, Lisbon, Portugal
2023*	Public Policy and the AI Revolution, Association for Public Policy Analysis & Management, Atlanta, GA
2023*	Department of Human Services, Allegheny County, Pittsburgh, PA
2023*	Quantitative Methods Workshop, Yale University, New Haven, CT
2023	Statistically Significant: Equity Concerns in Algorithmic Bias, Privacy, and Survey Representation, Joint Statistical Meetings, Toronto, CA
2023	K&L Gates Conference in Ethics and AI, Carnegie Mellon University, Pittsburgh, PA
2023	Multigroup Fairness and the Validity of Statistical Judgment, Simons Institute for the Theory of Computing, Berkeley, CA
2023	Automated Decision Systems Reading Group, University of California, Berkeley, CA
2023	Center for Information Technology Policy Lecture, Princeton University, Princeton, NJ
2023	Department of Computer Science, George Mason University, Fairfax, VA
2023	AI Seminar, New York University, New York, NY
2023	Data Science Initiative seminar, Brown University, Providence, RI
2023	Department of Engineering and Public Policy seminar, Carnegie Mellon University, Pittsburgh, PA
2023	Khoury College of Computer Sciences Lecture, Northeastern University, Boston, MA
2023	Department of Computer Science, University of Maryland, College Park, MD
2023	Halicioglu Data Science Institute and the School of Global Policy and Strategy, University of California, San Diego, CA
2023	Department of Statistics, University of California, Berkeley, CA
2023	McCourt School of Public Policy, Georgetown University, Washington, D.C.
2023	Information Science Colloquium Series, Cornell University, Ithaca, NY
2023	The Division of Decision, Risk, and Operations, Columbia Graduate School of Business, New York, NY
2023	The School of Data Science Colloquium, University of Virginia, Charlottesville, VA
2023	Econometrics & Statistics group, University of Chicago Booth School of Business, Chicago, IL
2022	Operations, Information, and Decisions Department, Wharton School of the University of Pennsylvania, Philadelphia, PA
2022	Symposium on Frontiers of Machine Learning & AI, University of Southern California, LA, CA
2022	INFORMS Session on Finding Sets of Near-Optimal Solutions for Mixed-Integer Programs, Indianapolis, IN
2022	American Mathematical Society Sectional Meeting on Causality, Amherst, MA (<i>declined</i>)
2022	Brown University Bravo Center Workshop on the Economics of Algorithms, Providence, RI
2022	Stanford University RegLab Summer Institute Speaker Series, Virtual
2021	Merck Data Science All Hands, Virtual
2021	Johns Hopkins University Causal Inference Working Group, Virtual

2021	PlaceKey COVID-19 Data Consortium, Virtual
2021	University of Pennsylvania Department of Biostatistics and Epidemiology, Virtual
2020	University of Chicago Crime Lab, Virtual

Doctoral Consortia

2022	EAAMO (ACM conference on Equity & Access in Algorithms, Mechanisms, and Optimization)
2022	FACCT (ACM Conference on Fairness, Accountability, and Transparency)
2020	FACCT (ACM Conference on Fairness, Accountability, and Transparency)
2019	AIES (AAAI / ACM Conference on Artificial Intelligence, Ethics, and Society)

Patents

2022	Enhancing Fairness in Transfer Learning for Machine Learning Models with Missing Protected Attributes in Source or Target Domains. Supriyo Chakraborty, <u>Amanda Coston</u> , Zairah Mustahsan, Karthikeyan Natesan Ramamurthy, Skyler Speakman, Kush R. Varshney, and Dennis Wei. US 11,443,236. <i>Granted</i> .
------	---

Service

Organization

2019-Now	Steering Committee of Machine Learning for Developing World (ML4D) NeurIPS Workshop
2019-2020	Co-organizer of Fairness, Ethics, Accountability, and Transparency Reading Group at CMU
2019-2020	Co-organizer of Machine Learning Department (MLD) Tea at CMU
2018-2019	Co-organizer of ML4D NeurIPS Workshop

Journal Referee

Nature Human Behaviour
 Journal of the Royal Statistical Society (JRSS-B)
 Journal of the American Statistical Association (JASA)
 Transactions on Machine Learning Research
 Data Mining and Knowledge Discovery

Program Committee and Conference Reviewer

2023	Ethical Reviewer, NeurIPS
2023	Program Committee, EAAMO
2023	Reviewer, ICLR
2022	Ethical Reviewer, NeurIPS
2022	Reviewer, NeurIPS
2022	Reviewer, NeurIPS Datasets and Benchmarks
2022	Program Committee, EAAMO
2022	Program Committee, FACCT
2022	Reviewer, ICML
2022	Reviewer, ICLR
2021	Area Chair, Responsible AI workshop at ICLR
2021	Ethical Reviewer, NeurIPS

2021	Reviewer, NeurIPS
2021	Reviewer, NeurIPS Datasets and Benchmarks
2021	Program Committee, FAccT
2021	Reviewer, ICML
2020	Reviewer, NeurIPS
2020	Program Committee, FAccT
2020	Reviewer, ICML
2020	Program Committee, AIES
2020	Program Committee, AAAI Emerging Track on AI for Social Impact
2019	Program Committee, IJCAI Workshop on AI for Social Good

Leadership

2012-2013	Committee on Discipline, Princeton University
2012-2013	Computer Science Undergraduate Council, Princeton University

Invited Conference & Workshop Roles

2022	Roundtable Lead for NeurIPS Workshop on Algorithmic Fairness through Lens of Causality
2022	Breakout Group Moderator for CCC & INFORMS Workshop II on AI/OR
2022	Breakout Group Moderator for NSF-Amazon Fairness in AI Principal Investigator meeting
2022	Session Chair for Responsible Data Management Session at FAccT

Teaching Experience

Teaching Assistant

2021 Spring	Introduction to Machine Learning (10-301/10-601), CARNEGIE MELLON UNIVERSITY
2012 Fall	Computers in our World (COS 109), PRINCETON UNIVERSITY

Project Instructor

2019 Summer	AI4ALL, CARNEGIE MELLON UNIVERSITY
	▷ Developed and led a project on algorithms, criminal justice, & fairness for high schoolers from historically excluded communities.

Mentorship

2022-Now	Women@SCS Mentor
2019-Now	CMU AI Mentor
2019	Women@SCS Roundtable Leader
2016-2017	Read Ahead Mentor
2014-2015	MySkills4Afrika (Microsoft) Virtual Mentor

Hackathon Distinctions

2015	Microsoft OneWeek Hackathon, Bing Finalist
	▷ Web answer to enable victims of revenge porn to remove content from Bing and OneDrive

- 2013 NYU-Abu Dhabi Hackathon for the Social Good, 2nd Place
 - ▷ Android app for sharing a travel route to facilitate safe travel for women
- 2012 Tiger Launch, Social Entrepreneurship, 3rd Place
 - ▷ Web service using QR codes to empower consumers to support value-aligned businesses

Civic Engagement

- 2014-2015 Court Appointed Special Advocate, Family Law CASA
 - ▷ Represented the child's interest in family law cases
- 2010-2012 Engineers Without Borders
 - ▷ Obtained & configured 50 One Laptop Per Child netbooks for a library in Ashaiman, Ghana
- 2007-2008 Congressional Intern, U.S. House of Representatives
 - ▷ Office of Congressman John Spratt representing South Carolina's 5th congressional district

Media Coverage

- 2021 "Smartphone Location Data Can Leave Out Those Most Hit by Covid-19." *Wall Street Journal*. <https://www.wsj.com/articles/...>
- 2020 "Stanford and Carnegie Mellon find race and age bias in mobility data that drives COVID-19 policy." *VentureBeat*. <https://venturebeat.com/ai/...>